

DATASTAX

DATASTAX

DATASTAX
ACADEMY




DataStax Enterprise

Foundations of Apache Cassandra™

DS201 - sections 3 - 5

July 27th, 2023

➤ #3 - Partitions



› Primary Keys

What is a Primary Key?

A table's primary key defines a unique path to access a single row of data.

› Primary Keys

Cassandra Primary Key Definition:

```
PRIMARY KEY ((partitionKey1, partitionKeyN),  
clusteringKey1, cluseringKeyN)
```

› Partition Keys

What is a Partition Key?

A partition key is the part of the Primary Key definition which tells Cassandra where to place the data in the cluster.

› Partition Keys

Cassandra Partition Key Definition:

```
PRIMARY KEY ((partitionKey1, partitionKeyN),  
clusteringKey1, cluseringKeyN)
```

› Table w/ a simple key

PRIMARY KEY (user_id)

```
> INSERT INTO users VALUES (2, 'Dev. Awesome', 'TX');
> INSERT INTO users VALUES (3, 'Epic Dev', 'NY');
> INSERT INTO users VALUES (4, 'IgotUr Data', 'TX');
> INSERT INTO users VALUES (5, 'Always Onomnom', 'TX');
> INSERT INTO users VALUES (6, 'Store Dat Data', 'NY');
> INSERT INTO users VALUES (7, 'Lolo Latency', 'CA');
> INSERT INTO users VALUES (8, 'Lovin Ur Bytes', 'TX');
> INSERT INTO users VALUES (9, 'Model De Tables', 'NY');
```

user_id	name	state
1	Dev Awesome	TX
2	Mrs. Reliable	CA
3	Epic Dev	NY
4	IgotUr Data	TX
5	Always Onomnom	TX
6	Store Dat Data	NY
7	Lolo Latency	CA
8	Luvn Ur Bytes	TX
9	Model Dat Data	NY

› Table w/ a simple key

PRIMARY KEY (user_id)

```
> SELECT user_id, state, name
   FROM users WHERE user_id=1;
```

```
user_id | state | name
-----+-----+-----
      1 |    TX | Dev Awesome
```

user_id	name	state
1	Dev Awesome	TX
2	Mrs. Reliable	CA
3	Epic Dev	NY
4	IgotUr Data	TX
5	Always Onomnom	TX
6	Store Dat Data	NY
7	Lolo Latency	CA
8	Luvn Ur Bytes	TX
9	Model Dat Data	NY

› However...

What if we want to query by
something other than user_id?

<shrug/>

Maybe we want to query by state.

› Partition by State

PRIMARY KEY
(state, user_id)

1	Dev Awesome	TX
2	Mrs. Reliable	CA
3	Epic Dev	NY
4	IgotUr Data	TX
5	Always Onomnom	TX
6	Store Dat Data	NY
7	Lolo Latency	CA
8	Luvn Ur Bytes	TX
9	Model Dat Data	NY

› Partition by State

PRIMARY KEY
(state, user_id)

TX



1	Dev Awesome	TX
4	IgotUr Data	TX
5	Always Onomnom	TX
8	Luvn Ur Bytes	TX

NY



3	Epic Dev	NY
6	Store Dat Data	NY
9	Model Dat Data	NY

CA



2	Mrs. Reliable	CA
7	Lolo Latency	CA

Murmur3
Partitioner

› Murmur3 Partitioner

PRIMARY KEY
(state, user_id)

-2594951604484898973

7836943455311205863

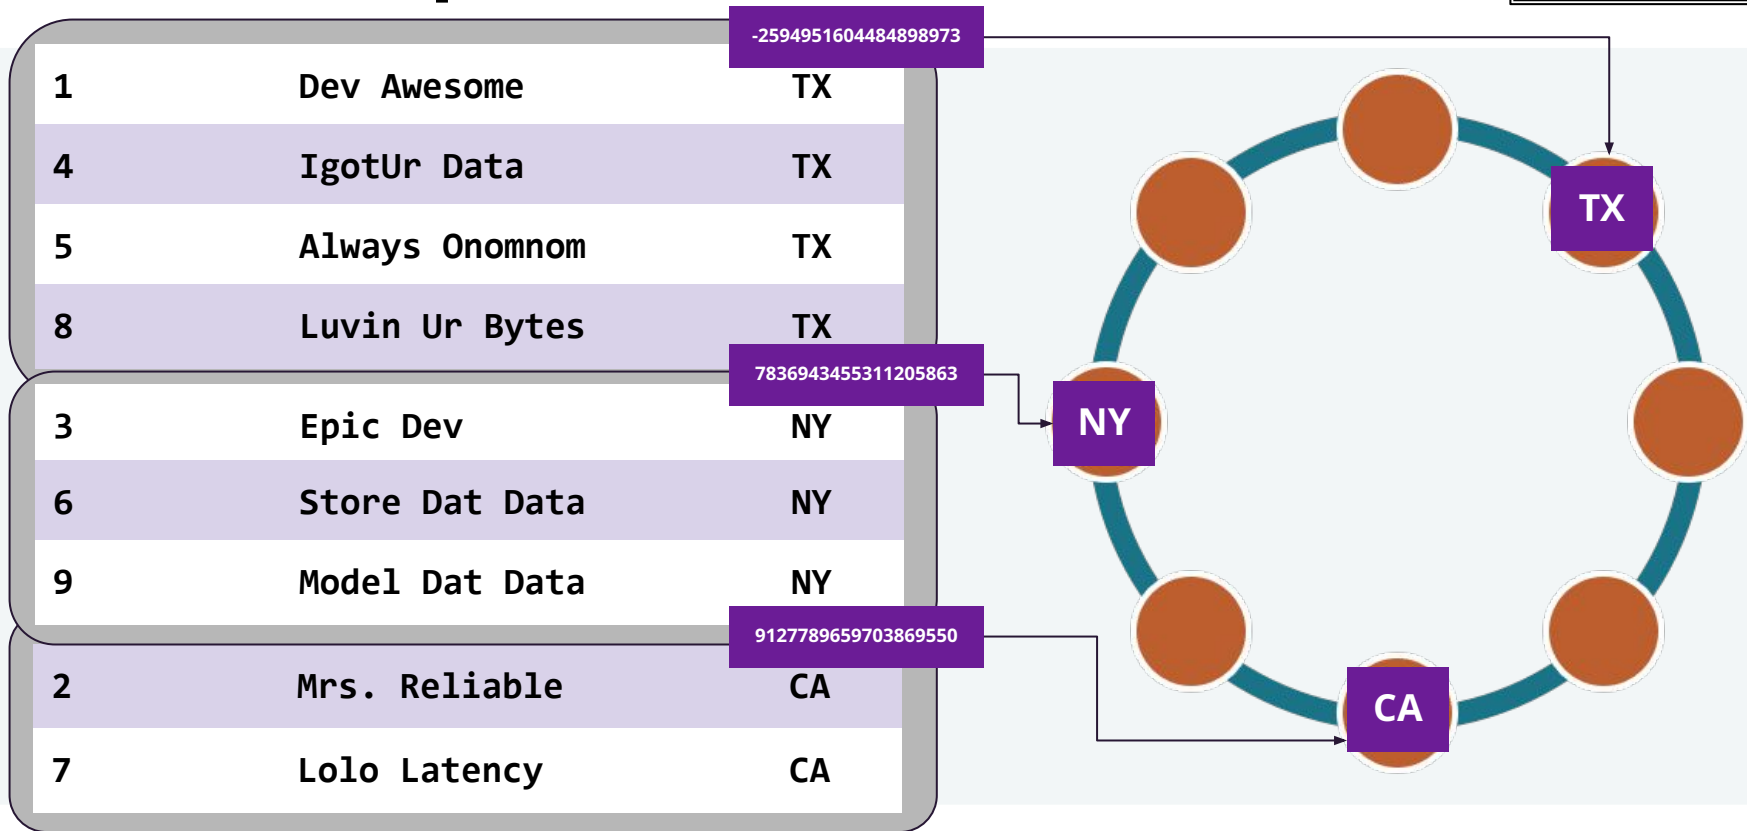
9127789659703869550

1	Dev Awesome	TX
4	IgotUr Data	TX
5	Always Onomnom	TX
8	Luvn Ur Bytes	TX

3	Epic Dev	NY
6	Store Dat Data	NY
9	Model Dat Data	NY

2	Mrs. Reliable	CA
7	Lolo Latency	CA

› Data is placed in the cluster




➤ Exercise #3



Hands-on Exercise #3

- › Create tables with different partition keys
- › Execute CQL queries on tables with different partition keys

➤ #4 - Clustering Columns



› Clustering Columns

What is a Clustering Column?

A clustering column is the part of the Primary Key definition which defines the on-disk sort order within a partition.

› Primary Keys

Cassandra Clustering Column Definition:

```
PRIMARY KEY ((partitionKey1, partitionKeyN),  
clusteringKey1, cluseringKeyN)
```

› Effects of different PK definitions

PRIMARY KEY (user_id)

*Assume the rows are written
in the order of their user_id.*

5	Always Onomnom	TX	Dallas
1	Dev Awesome	TX	Houston
8	Luvn Ur Bytes	TX	Santa Fe
2	ComeTo Learnin	TX	Dallas
4	IgotUr Data	TX	Austin
7	Data Rowman	TX	Austin
6	Lone Star	TX	El Paso
9	Compact One	TX	Houston
3	Lone Node	TX	Snyder

› Effects of different PK definitions

PRIMARY KEY ((state), city)

*Assume the rows are written
in the order of their user_id.*

4	IgotUr Data	TX	Austin
7	Data Rowman	TX	Austin
2	ComeTo Learnin	TX	Dallas
5	Always Onomnom	TX	Dallas
6	Lone Star	TX	El Paso
1	Dev Awesome	TX	Houston
9	Compact One	TX	Houston
8	Luvn Ur Bytes	TX	Santa Fe
3	Lone Node	TX	Snyder

Tips

› From the Trenches

Cassandra Primary Keys are unique!

For multiple writes to the same PK, *the last write wins!*

› Effects of different PK definitions

PRIMARY KEY
((state), city, user_id))

-2594951604484898973

4	IgotUr Data	TX	Austin
7	Data Rowman	TX	Austin
2	ComeTo Learnin	TX	Dallas
5	Always Onomnom	TX	Dallas
6	Lone Star	TX	El Paso
1	Dev Awesome	TX	Houston
9	Compact One	TX	Houston
8	Luvn Ur Bytes	TX	Santa Fe
3	Lone Node	TX	Snyder

› Clustering Columns

Querying

- You must first provide the (complete) partition key.
- Clustering columns must be specified *in order*.
- You can perform equality (=) or range queries (< , <=) on clustering columns.
- Equality comparisons must be performed before range comparisons.
- Since data is sorted on disk, range queries are a binary search followed by a sequential read.

› Clustering Column queries

Query:

```
> SELECT state, city, user_id, name  
FROM users_by_city_by_state  
WHERE state= 'TX' AND city='Austin';
```

state	city	user_id	name
TX	Austin	4	IgotUr Data
TX	Austin	7	Data Rowman

› Clustering Column queries

Query:

```
> SELECT state, city, user_id, name  
FROM users_by_city_by_state  
WHERE state= 'TX' AND city='Austin'  
AND user_id > 5;
```

state	city	user_id	name
TX	Austin	7	Data Rowman

› Clustering Column queries

Query:

```
> SELECT state, city, user_id, name  
FROM users_by_city_by_state  
WHERE state= 'TX' AND user_id > 5;
```

```
[Invalid query] message="PRIMARY KEY column  
"user_id" cannot be restricted as preceding  
column "city" is not restricted"
```

› Changing Default Ordering

Clustering Column definitions

- Default to ASCending sort direction.
- Sort direction can be influenced by the **WITH CLUSTERING ORDER BY** clause.
- Must include up to and including the columns you wish to order DESCending.

Tips

› From the Trenches

Queries filtered by clustering columns *perform faster* if you do not switch the table's sort order!

Sometimes you may want to duplicate a table with a different sort order.

› Clustering Column queries

Error:

[Invalid query] message="Cannot execute this query as it might involve data filtering and thus may have unpredictable performance. If you want to execute this query despite the performance unpredictability, use ALLOW FILTERING."

› Clustering Column queries

Allow Filtering

- Relaxes the query requirements on partition keys.
- You can then query on just clustering columns.
- Causes Cassandra to scan all partitions in a table.
- DO NOT USE IT:
 - Might be ok on small data sets or small clusters.
 - But seriously just don't.

**Pop
Quiz!**

**When should
you use
ALLOW
FILTERING?**



Never

› Guardrails in Cassandra 4.1

cassandra.yaml:

```
# Guardrail to allow/disallow querying  
# with ALLOW FILTERING. Defaults to true.  
  
allow_filtering_enabled: false
```

➤ Exercise #4



Hands-on Exercise #4

- › Create tables with different clustering columns
- › Execute CQL queries on tables with different clustering columns

➤ #5 - Application Connectivity



› Drivers

There are drivers for most languages:

- Java
- Python
- C#
- C++
- Node.js
- Ruby
- Go

› Similar Across Most Languages

Python:

```
auth = PlainTextAuthProvider(username=user,  
                             password=pwd)  
cluster = Cluster(["127.0.0.1"], auth_provider=auth)  
session = cluster.connect("killrvideo")  
result = session.execute("<query>")[0]
```

› Similar Across Most Languages

Java:

```
CqlSession session= CqlSession.builder()  
    .withKeyspace("killrvideo")  
    .setAuthCredentials(user, pwd)  
    .build();  
ResultSet results = session.execute("<query>");
```


› Similar Across Most Languages

C#:

```
var cluster = Cluster.Builder()  
    .AddContactPoint("127.0.0.1")  
    .WithAuthProvider(  
        new PlainTextAuthProvider(user, pwd))  
    .Build();  
var session = cluster.Connect("killrvideo");  
var rs = session.Execute("<query>");
```

› Previous Workshop

Cassandra Application Development

[https://github.com/datastaxdevs/
workshop-cassandra-application-development](https://github.com/datastaxdevs/workshop-cassandra-application-development)



<https://www.youtube.com/watch?v=nPps2oqWhnY>

› Setup

Python

```
$ pip install cassandra-driver
```

```
from cassandra.cluster import Cluster  
from cassandra.auth import PlainTextAuthProvider  
import sys
```

› Parameters

Set defaults

```
hostname="127.0.0.1"  
username="cassandra"  
password="cassandra"  
protocol=4
```

```
if (len(sys.argv) > 1):  
    hostname=sys.argv[1]  
if (len(sys.argv) > 2):  
    username=sys.argv[2]  
if (len(sys.argv) > 3):  
    password=sys.argv[3]
```

› Build connection

```
nodes = []  
nodes.append(hostname)  
auth_provider =  
    PlainTextAuthProvider(username=username,  
                           password=password)  
cluster = Cluster(nodes, auth_provider=auth_provider,  
                  protocol_version=protocol)  
session = cluster.connect()
```

› Query Data

SELECT

```
rows = session.execute(  
    "SELECT broadcast_address, host_id,  
    data_center, rack, tokens  
    FROM system.local; -- this is a comment\n")
```

› Display Results

Iterate & print

```
for row in rows:
```

```
    print("address: " + row[0])
```

```
    print("host id: " + str(row[1]))
```

```
    print("data center: " + row[2])
```

```
    print("rack: " + row[3])
```

```
    print("tokens: " + str(len(row[4])))
```

➤ No Exercises for section #5