

DATASTAX

DATASTAX

DATASTAX

ACADEMY




# DataStax Enterprise

**Java Development with Apache Cassandra™**

**July 27th, 2023**

# ➤ #1 – Introduction





# Learning Objectives

- › Building a Connection
- › Simple Queries
- › Simple Writes
- › Prepared Statements
- › BATCH updates

# › Maven Dependencies

**pom.xml:**

```
<dependency>  
  <groupId>com.datastax.oss</groupId>  
  <artifactId>java-driver-core</artifactId>  
  <version>4.16.0</version>  
</dependency>
```

## › Maven Dependencies

**pom.xml:**

```
<dependency>  
  <groupId>com.datastax.astra</groupId>  
  <artifactId>astra-spring-boot-starter</artifactId>  
  <version>0.6</version>  
</dependency>
```

## › Building a connection

### Java:

```
CqlSession session= CqlSession.builder()  
    .appContactPoints(endpointList)  
    .withKeyspace(keyspace)  
    .setAuthCredentials(user, pwd)  
    .withLocalDatacenter(datacenterName)  
    .build();
```

## › Cassandra Connection

### Tips:

- Build and reuse one session object.
- Gracefully shutdown using `finalize()`.
- Pass-in credentials as environment variables.
- Specify a default data center.
- Specify a default keyspace.



## › Simple Queries

### Java:

```
String strCQL = "SELECT * FROM system.local LIMIT 1";  
ResultSet rs = session.execute(strCQL);  
Row localInfo = rs.one();  
  
String address = localInfo.getString("rpc_address");  
UUID hostId = localInfo.getUuid("host_id");  
int port = localInfo.getInt("rpc_port");
```

## › Prepared Statement Queries

### Java:

```
String strCQL = "SELECT name, brand FROM product WHERE  
product_group = ?";  
PreparedStatement prepared = session.prepare(strCQL);  
BoundStatement bound = prepared.bind(productGroup);  
ResultSet rs = session.execute(bound);  
List<Row> products = rs.all();
```

## › Vector Search

### Java:

```
PreparedStatement qvPrep = session.prepare(  
    "SELECT * FROM product_vector ORDER BY  
    product_vector ANN OF ? LIMIT 2;");  
BoundStatement qvBound = qvPrep.bind(prd.getVector());  
ResultSet rsV = session.execute(qvBound);  
List<Row> ann = rsV.all();
```

# › Cassandra Queries

## Tips:

- Name the columns in the SELECT clause.
- ALWAYS use a WHERE or a LIMIT clause.
- Never use ALLOW FILTERING.
- For repeated queries, use a prepared statement:
  - The prepared statement cache is your friend!
  - Prepare *outside* of the loop.
  - Execute *inside* of the loop.

## › Prepared Statement Writes

Java:

```
String strCQL = "INSERT INTO product (product_id,  
    name, brand) VALUES (?, ?, ?)";  
  
PreparedStatement prepared = session.prepare(strCQL);  
  
BoundStatement bound = prepared.bind(productId,  
    name, brand);  
  
session.execute(bound);
```

## › Cassandra Batch – correct use

### Java:

```
UUID productId = UUID.randomUUID();
```

```
BoundStatement productBoundStatement1 = this.getProductInsertStatement(product, productId,  
PRODUCT_BY_ID_TABLE_NAME);
```

```
BoundStatement productBoundStatement2 = this.getProductInsertStatement(product, productId,  
PRODUCT_BY_NAME_TABLE_NAME);
```

```
BatchStatement batch = BatchStatement.newInstance(DefaultBatchType.LOGGED,  
productBoundStatement1,productBoundStatement2);
```

```
session.execute(batch);
```

# › Cassandra Batch – incorrect use

## CQL:

### BEGIN BATCH

```
INSERT INTO cycling.cyclist_name (id, lastname, firstname) VALUES  
(6d5f1663-89c0-45fc-8cfd-60a373b01622, 'HOSKINS', 'Melissa');
```

```
INSERT INTO cycling.cyclist_name (id, lastname, firstname) VALUES  
(38ab64b6-26cc-4de9-ab28-c257cf011659, 'FERNANDES', 'Marcia');
```

```
INSERT INTO cycling.cyclist_name (id, lastname, firstname) VALUES  
(9011d3be-d35c-4a8d-83f7-a3c543789ee7, 'NIEWIADOMA', 'Katarzyna');
```

```
INSERT INTO cycling.cyclist_name (id, lastname, firstname) VALUES  
(95addc4c-459e-4ed7-b4b5-472f19a67995, 'ADRIAN', 'Vera');
```

### APPLY BATCH;

## › Cassandra Writes

### Tips:

- Don't INSERT nulls! Those are “tombstones.”
- Don't run in-place updates.
- Don't run many deletes.
- BATCH for one update across multiple tables - Good!
- BATCH for multiple updates across one table - Bad!



# › Cassandra Bulk Writes – correct

## Java:

```
for (BoundStatement boundStatement : cqlInserts) {  
    ResultSetFuture future = session.executeAsync(boundStatement);  
    futures.add(future);  
    threadCount++;  
    if (threadCount > 19) {  
        futures.forEach(ResultSetFuture::getUninterruptibly);  
        futures = new ArrayList<>();  
        threadCount = 0;  
    }  
}
```

Iterate through a  
**List<BoundStatement>**  
Execute each  
Add to  
Statement  
**List<ResultSetFuture>**

If too many threads  
are in-flight, wait to  
ensure completion,  
restart threadCount  
and futures.

# › Java Exercise



# Hands-on Exercise Java

- › Build a simple, restful product/recommendation service
- › Support query by product\_id
- › Support query by vector ANN

## › Exercises



<https://github.com/aar0np/DS-Java-DSE>